

academicPerformance

May 11, 2022

0.0.1 Data Wrangling II

Create an “Academic performance” dataset of students and perform the following operations using Python. - Scan all variables for missing values and inconsistencies. If there are missing values and/or inconsistencies, use any of the suitable techniques to deal with them. - Scan all numeric variables for outliers. If there are outliers, use any of the suitable techniques to deal with them. - Apply data transformations on at least one of the variables. The purpose of this transformation should be one of the following reasons: to change the scale for better understanding of the variable, to convert a non-linear relation into a linear one, or to decrease the skewness and convert the distribution into a normal distribution. Reason and document your approach properly.

```
[76]: #import libraries and dataset
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
#Dataset CSV
url = "eduData.csv"
df = pd.read_csv(url)
print(df.head(10))
```

	gender	NationalITY	PlaceofBirth	StageID	GradeID	SectionID	Topic	\
0	NaN	KW	KuwaIT	lowerlevel	G-04	A	IT	
1	M	KW	NaN	lowerlevel	G-04	A	NaN	
2	M	KW	KuwaIT	NaN	G-04	A	IT	
3	M	KW	KuwaIT	lowerlevel	G-04	A	IT	
4	NaN	KW	KuwaIT	lowerlevel	G-04	A	IT	
5	F	KW	KuwaIT	lowerlevel	G-04	A	IT	
6	M	KW	KuwaIT	MiddleSchool	G-07	A	NaN	
7	M	KW	NaN	MiddleSchool	G-07	A	Math	
8	F	KW	KuwaIT	MiddleSchool	G-07	A	Math	
9	F	KW	KuwaIT	MiddleSchool	G-07	B	IT	

	Semester	Relation	cns	dsa	oops	os
0	F	Father	NaN	16.0	2	20
1	F	Father	20.0	20.0	3	25
2	F	Father	10.0	7.0	0	30
3	F	Father	NaN	25.0	5	35
4	F	Father	40.0	50.0	12	50

5	F	Father	42.0	30.0	13	70
6	F	Father	35.0	12.0	0	17
7	F	NaN	NaN	NaN	15	22
8	F	Father	12.0	21.0	16	50
9	F	Father	NaN	80.0	25	70

```
[77]: #check no null value in each column
print(df.isnull().sum())
```

```
gender          6
NationalITy     1
PlaceofBirth    5
StageID         2
GradeID         1
SectionID       0
Topic          4
Semester        0
Relation        2
cns             7
dsa             1
oops            0
os              0
dtype: int64
```

0.0.2 Ways to fill the null values

- simply drop the row having null value
- imputate with mean, median or mode.
- fill with random value like “Unknown”
- replace categorical variable with previous value
- replace inconsitent data with null value.
- imputate by interpolation

```
[79]: #drop the whole row which is having NULL value
t=df.dropna()
print(t.isnull().sum())
print("Before dropping null values:- ",t.shape)
print("After dropping null values:- ",df.shape)
```

```
gender          0
NationalITy     0
PlaceofBirth    0
StageID         0
GradeID         0
SectionID       0
Topic          0
Semester        0
Relation        0
```

```
cns          0
dsa          0
oops         0
os           0
dtype: int64
Before dropping null values:- (9, 13)
After dropping null values:- (28, 13)
```

```
[80]: #imputation by mean
url = "eduData.csv"
df = pd.read_csv(url)
df["cns"] = df["cns"].replace(np.NaN, df["cns"].mean())

print(df["cns"])
```

```
0    25.571429
1    20.000000
2    10.000000
3    25.571429
4    40.000000
5    42.000000
6    35.000000
7    25.571429
8    12.000000
9    25.571429
10   50.000000
11   19.000000
12    5.000000
13   20.000000
14   25.571429
15   30.000000
16   36.000000
17   25.571429
18   69.000000
19   70.000000
20   25.571429
21   10.000000
22   15.000000
23    2.000000
24    0.000000
25    8.000000
26   19.000000
27   25.000000
```

```
Name: cns, dtype: float64
```

imputation using interpolation Linear Interpolation simply means to estimate a missing value by connecting dots in a straight line in increasing order. In short, It estimates the unknown value in the same increasing order from previous values.

```
[82]: import statistics
df = pd.read_csv(url)
df["cns"] = df["cns"].interpolate(method='linear')
print(df["cns"])
```

```
0      NaN
1      20.0
2      10.0
3      25.0
4      40.0
5      42.0
6      35.0
7      23.5
8      12.0
9      31.0
10     50.0
11     19.0
12      5.0
13     20.0
14     25.0
15     30.0
16     36.0
17     52.5
18     69.0
19     70.0
20     40.0
21     10.0
22     15.0
23      2.0
24      0.0
25      8.0
26     19.0
27     25.0
Name: cns, dtype: float64
```

```
[83]: #replace categorical variable with random value
df["gender"] = df["gender"].fillna('unknown')
print(df["gender"])
```

```
0      unknown
1           M
2           M
3           M
4      unknown
5           F
6           M
7           M
8           F
```

```

9          F
10         M
11         M
12         M
13         M
14         F
15         F
16    unknown
17         M
18         F
19    unknown
20         F
21         F
22         M
23    unknown
24         M
25         M
26    unknown
27         M
Name: gender, dtype: object

```

```

[84]: #replace categorical variable with previous value
df = pd.read_csv(url)
df["gender"] = df["gender"].fillna(method='ffill')
print(df["gender"])

```

```

0      NaN
1       M
2       M
3       M
4       M
5       F
6       M
7       M
8       F
9       F
10      M
11      M
12      M
13      M
14      F
15      F
16      F
17      M
18      F
19      F
20      F
21      F

```

```
22      M
23      M
24      M
25      M
26      M
27      M
Name: gender, dtype: object
```

```
[85]: df = pd.read_csv(url)
      #creating the inconsistent data
      df["gender"]=df["gender"].fillna(100)
      cnt=0
      for row in df["gender"]:
          try:
              int(row)
              df.loc[cnt,"gender"]=np.nan
          except ValueError:
              pass
          cnt+=1

      print(df["gender"])
```

```
0      NaN
1       M
2       M
3       M
4      NaN
5       F
6       M
7       M
8       F
9       F
10      M
11      M
12      M
13      M
14      F
15      F
16      NaN
17      M
18      F
19      NaN
20      F
21      F
22      M
23      NaN
24      M
25      M
```


```
26     NaN
27     M
Name: gender, dtype: object
```

```
[86]: #data Tranformation to decrease the skewness
# Skewness is a measure of the asymmetry of the probability distribution of a
#real-valued random variable about its mean.
df.skew(numeric_only=True)
```

```
[86]: cns      0.946321
dsa      1.241056
oops     0.440346
os       0.056839
dtype: float64
```

```
[87]: df.hist(grid=False,
          figsize=(10, 6),
          bins=30)
```

```
[87]: array([[<AxesSubplot:title={'center':'cns'}>,
          <AxesSubplot:title={'center':'dsa'}>],
        [<AxesSubplot:title={'center':'oops'}>,
          <AxesSubplot:title={'center':'os'}>]], dtype=object)
```




```
[88]: df.insert(len(df.columns), 'dsa_Sqrt',  
              np.sqrt(df.dsa))
```

```
[89]: df.skew(numeric_only=True)
```

```
[89]: cns      0.946321  
dsa      1.241056  
oops     0.440346  
os       0.056839  
dsa_Sqrt  0.291450  
dtype: float64
```


```
[90]: #identify outliers and handle them  
sns.boxplot(x=df["os"])
```

```
[90]: <AxesSubplot:xlabel='os'>
```




```
[91]: sns.boxplot(x=df["cns"])
```

```
[91]: <AxesSubplot:xlabel='cns'>
```

```
[92]: #We can clearly see that 3 values greater than 60 are outliers.
sns.boxplot(x=df["dsa"])
```

```
[92]: <AxesSubplot:xlabel='dsa'>
```



```
[93]: print(np.where(df['dsa']>65))
outliers=np.where(df['dsa']>65)
df.head(10)
df.shape
```

```
(array([ 9, 10, 14]),)
```

```
[93]: (28, 14)
```

```
[94]: print(df)
```

	gender	NationalITy	PlaceofBirth	StageID	GradeID	SectionID	Topic \
0	NaN	KW	KuwaIT	lowerlevel	G-04	A	IT
1	M	KW	NaN	lowerlevel	G-04	A	NaN
2	M	KW	KuwaIT	NaN	G-04	A	IT
3	M	KW	KuwaIT	lowerlevel	G-04	A	IT
4	NaN	KW	KuwaIT	lowerlevel	G-04	A	IT
5	F	KW	KuwaIT	lowerlevel	G-04	A	IT
6	M	KW	KuwaIT	MiddleSchool	G-07	A	NaN
7	M	KW	NaN	MiddleSchool	G-07	A	Math
8	F	KW	KuwaIT	MiddleSchool	G-07	A	Math
9	F	KW	KuwaIT	MiddleSchool	G-07	B	IT
10	M	KW	KuwaIT	MiddleSchool	G-07	A	Math
11	M	KW	KuwaIT	MiddleSchool	G-07	B	Math
12	M	KW	KuwaIT	lowerlevel	NaN	A	IT
13	M	lebanon	lebanon	NaN	G-08	A	Math
14	F	KW	KuwaIT	MiddleSchool	G-08	A	Math
15	F	KW	KuwaIT	MiddleSchool	G-06	A	IT
16	NaN	NaN	KuwaIT	MiddleSchool	G-07	B	IT
17	M	KW	NaN	MiddleSchool	G-07	A	NaN
18	F	KW	KuwaIT	MiddleSchool	G-07	A	IT
19	NaN	KW	KuwaIT	MiddleSchool	G-07	B	IT
20	F	KW	NaN	MiddleSchool	G-07	A	IT
21	F	KW	KuwaIT	MiddleSchool	G-07	B	IT
22	M	KW	KuwaIT	MiddleSchool	G-07	A	IT
23	NaN	KW	KuwaIT	MiddleSchool	G-07	A	IT
24	M	KW	KuwaIT	MiddleSchool	G-07	B	NaN
25	M	KW	NaN	MiddleSchool	G-07	A	IT
26	NaN	KW	KuwaIT	MiddleSchool	G-07	B	IT
27	M	KW	KuwaIT	MiddleSchool	G-08	A	Arabic

	Semester	Relation	cns	dsa	oops	os	dsa_Sqrt
0	F	Father	NaN	16.0	2	20	4.000000
1	F	Father	20.0	20.0	3	25	4.472136
2	F	Father	10.0	7.0	0	30	2.645751

3	F	Father	NaN	25.0	5	35	5.000000
4	F	Father	40.0	50.0	12	50	7.071068
5	F	Father	42.0	30.0	13	70	5.477226
6	F	Father	35.0	12.0	0	17	3.464102
7	F	NaN	NaN	NaN	15	22	NaN
8	F	Father	12.0	21.0	16	50	4.582576
9	F	Father	NaN	80.0	25	70	8.944272
10	F	Father	50.0	88.0	30	80	9.380832
11	F	Father	19.0	6.0	19	12	2.449490
12	F	Father	5.0	1.0	0	11	1.000000
13	F	Father	20.0	14.0	12	19	3.741657
14	F	NaN	NaN	70.0	44	60	8.366600
15	F	Father	30.0	40.0	22	66	6.324555
16	F	Father	36.0	30.0	20	80	5.477226
17	F	Father	NaN	13.0	35	90	3.605551
18	F	Mum	69.0	15.0	36	96	3.872983
19	F	Mum	70.0	50.0	40	99	7.071068
20	F	Father	NaN	60.0	33	90	7.745967
21	F	Father	10.0	12.0	4	80	3.464102
22	F	Father	15.0	21.0	2	90	4.582576
23	F	Father	2.0	0.0	2	50	0.000000
24	F	Father	0.0	2.0	3	70	1.414214
25	F	Father	8.0	7.0	30	40	2.645751
26	F	Father	19.0	19.0	25	40	4.358899
27	F	Father	25.0	15.0	12	33	3.872983

```
[96]: new_df = df.drop(df.index[outliers])
```

```
[101]: print("Previous size :-",df.shape)
        print("Current size :- ",new_df.shape)
```

```
Previous size :- (28, 14)
```

```
Current size :- (25, 14)
```